

CS251 Fall 2020  
(cs251.stanford.edu)



# Cryptocurrencies and Blockchain Technologies

Dan Boneh

Benedikt Bünz

Stanford University

# What is a blockchain?

Abstract answer: a blockchain provides

- coordination between many parties,
- when there is no single trusted party

if trusted party exists  $\Rightarrow$  no need for a blockchain

[financial systems: often no trusted party]

# What is all the excitement about?

(1) Basic application: a digital currency (stored value)

- Current largest: Bitcoin (2009), Ethereum (2015), Tether (2014)
  - Global: accessible to anyone with an Internet connection

Opinion

The New York Times

# Bitcoin Has Saved My Family

“Borderless money” is more than a buzzword when you live in a collapsing economy and a collapsing dictatorship.

# What is all the excitement about?

(1) Basic application: a digital currency (stored value)

- Current largest: Bitcoin (2009), Ethereum (2015), Tether (2014)
  - Global: accessible to anyone with an Internet connection
- 

(2) Beyond stored value: **decentralized applications (DAPPs)**

- DeFi: financial applications managed by public programs
    - examples: stablecoins, lending, exchanges, ....
  - Asset management (e.g., domain names, games).
- 

(3) New programming model: writing decentralized programs

# Transaction volume

24h Volume

(Sep, 2020)

 Bitcoin \$70,163,302,153

 Ethereum \$62,307,903,847

 Tether \$52,715,790,830

 XRP \$1,724,384,881

# Central Bank Digital Currency (CBDC)



By Thomas Simms

**F**orbes.com  
China's Digital Currency Is Ready,  
**C**entral Bank Says  
30 CENTS  
PUBLISHED WED, OCT 16, 2019  
BY THOMAS SIMMS  
PHOTOGRAPH BY CHINA DAILY/REUTERS  
AUG 11, 2019

on retail CBDC

2019

# What is a blockchain?

Layer 3:

**user facing tools** (cloud servers)

Layer 2:

**applications** (DAPPs, smart contracts)

Layer 1.5:

**compute layer** (blockchain computer)

Layer 1:

**consensus layer**

# Consensus layer (informal)

A public append-only data structure:

achieved by replication

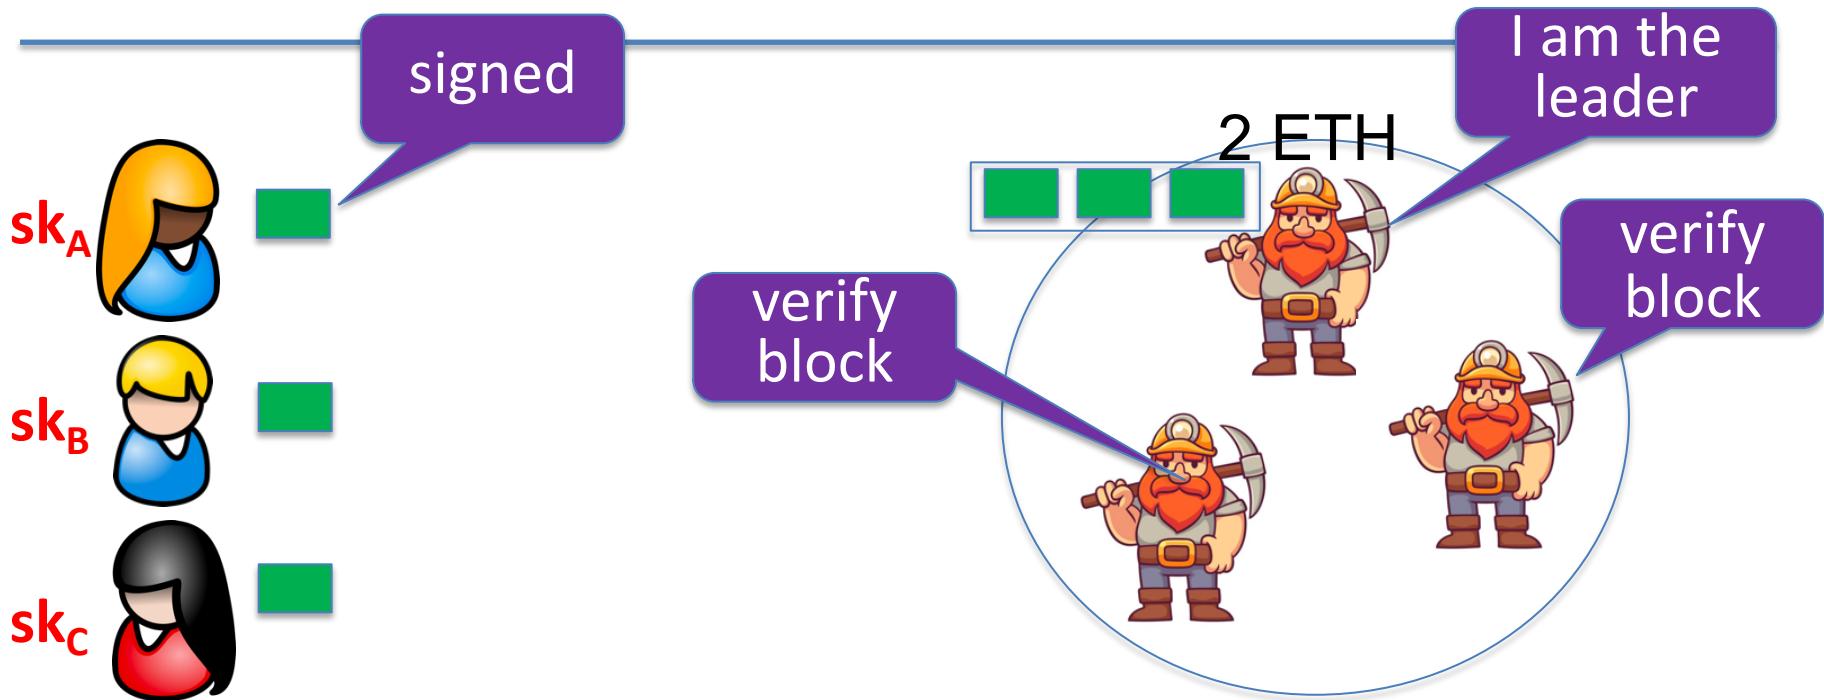
- **Persistence:** once added, data can never be removed\*
- **Consensus:** all honest participants have the same data\*\*
- **Liveness:** honest participants can add new transactions
- **Open(?)**: anyone can add data (no authentication)

Layer 1:

consensus layer

# How are blocks added to chain?

blockchain

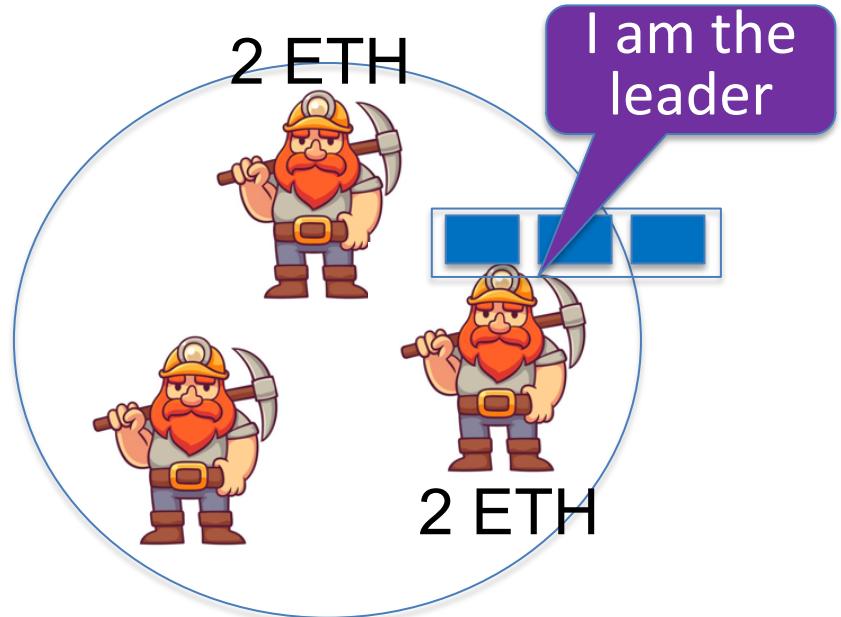
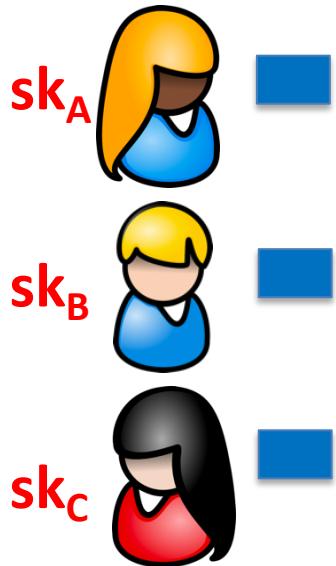


# How are blocks added to chain?

blockchain



...



# Layer 1.5: The blockchain computer

**DAPP logic is encoded in a program that runs on blockchain**

- Rules are enforced by a public program (public source code)  
⇒ transparency: no single trusted 3<sup>rd</sup> party
- The DAPP program is executed by parties who create new blocks  
⇒ public verifiability: everyone can verify state transitions

Layer 1.5:

compute layer

Layer 1:

consensus layer

# Layer 2: Decentralized applications (DAPPS)

Run on  
blockchain  
computer



Layer 2: **applications** (DAPPs, smart contracts)

Layer 1.5: blockchain computer

Layer 1: consensus layer

# Layer 3: Common DAPP architecture

Layer 3: user facing servers



end user

DAPP

DAPP

DAPP



on-chain  
state

blockchain computer

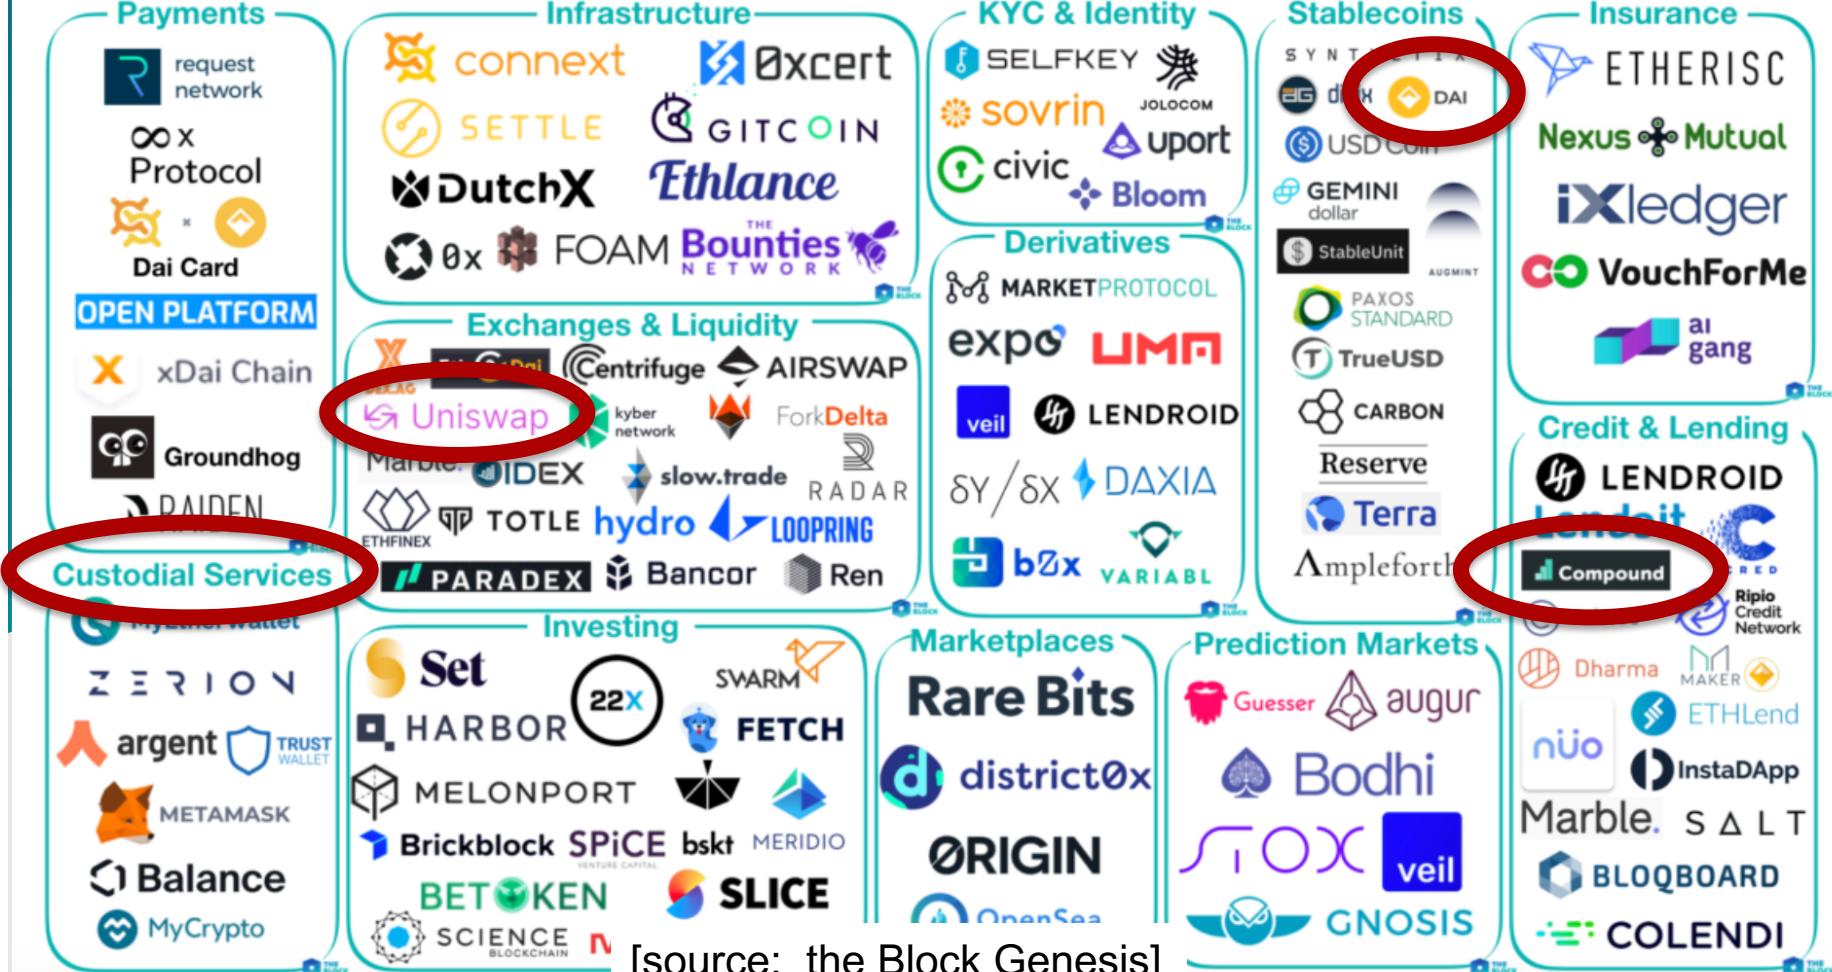
(layer 2)

(layer 1.5)

(layer 1)

consensus layer

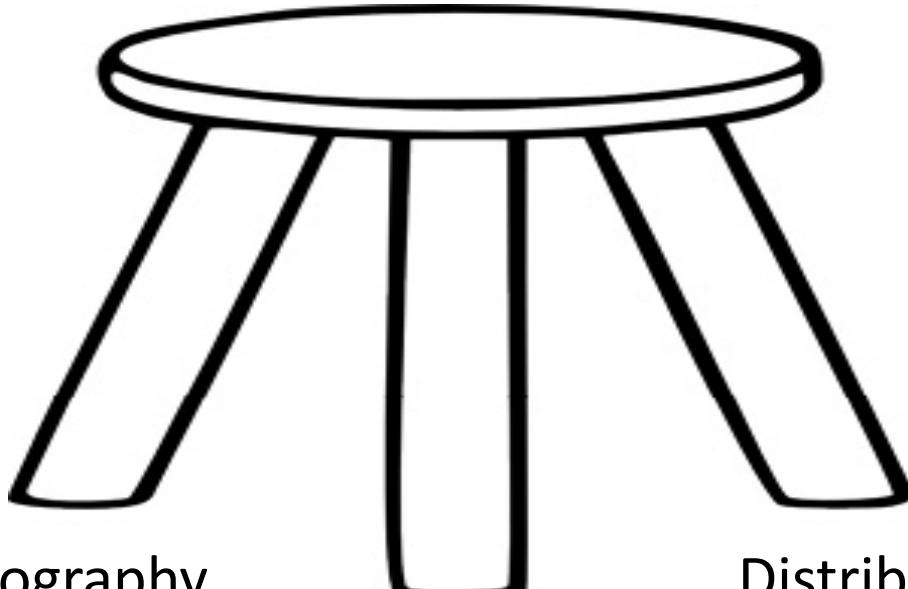
# Ethereum's DeFi



# lots of experiments ...

DEFI PULSE	Name	Chain	Category	Locked (USD) ▼
🏆 1.	Aave	Ethereum	Lending	\$1.49B
🥈 2.	Maker	Ethereum	Lending	\$1.26B
🥉 3.	Curve Finance	Ethereum	DEXes	\$1.00B
4.	yearn.finance	Ethereum	Assets	\$785.8M
5.	Synthetix	Ethereum	Derivatives	\$769.4M
6.	Compound	Ethereum	Lending	\$626.5M
7.	WBTC	Ethereum	Assets	\$570.7M
8.	Uniswap	Ethereum	DEXes	\$564.5M

# This course



Cryptography

Economics

Distributed systems

# Course organization

1. The starting point: Bitcoin mechanics
2. Consensus protocols
3. Ethereum and decentralized applications
4. Economics of decentralized applications
5. Scaling the blockchain: 10K Tx/sec and more
6. Private transactions on a public blockchain  
(SNARKs and zero knowledge proofs)

# Course organization

cs251.stanford.edu

- Three homework problems, four projects, final exam(?)
- Optional weekly sections on Friday

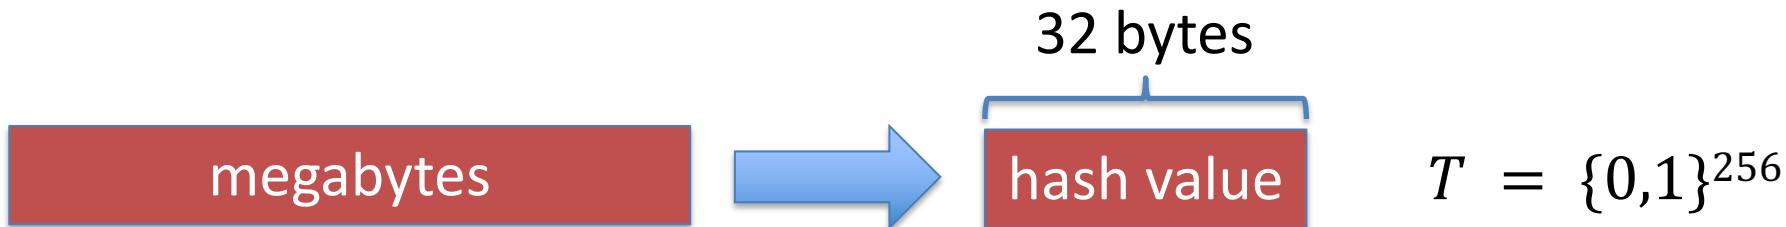
Please tell us how we can improve ...  
Don't wait until the end of the quarter

Let's get started ...

# Cryptography Background

## (1) cryptographic hash functions

An efficiently computable function  $H: M \rightarrow T$   
where  $|M| \gg |T|$



# Collision resistance

Def: a collision for  $H: M \rightarrow T$  is pair  $x \neq y \in M$  s.t.

$$H(x) = H(y)$$

$|M| \gg |T|$  implies that many collisions exist

Def: a function  $H: M \rightarrow T$  is collision resistant if it is “hard” to find even a single collision for  $H$  (we say  $H$  is a CRHF)

Example: **SHA256:**  $\{x : \text{len}(x) < 2^{64} \text{ bytes}\} \rightarrow \{0,1\}^{256}$

details in CS255

# An application: committing to data

Alice has a large file  $m$ . She publishes  $h = H(m)$  (32 bytes)

Bob has  $h$ . Later he learns  $m'$  s.t.  $H(m') = h$

$H$  is a CRHF  $\Rightarrow$  Bob is convinced that  $m' = m$

(otherwise,  $m$  and  $m'$  are a collision for  $H$ )

We say that  $h = H(m)$  is a **binding commitment** to  $m$

(note: not hiding,  $h$  may leak information about  $m$ )

# Committing to a list (of transactions)

Alice has  $S = (m_1, m_2, \dots, m_n)$

32 bytes

Goal:

- Alice publishes a short binding commitment to  $S$ ,  $h = \text{commit}(S)$
- Bob has  $h$ . Given  $(m_i, \text{proof } \pi_i)$  can check that  $S[i] = m_i$

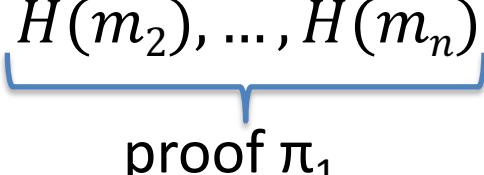
Bob runs  $\text{verify}(h, i, m_i, \pi_i) \rightarrow \text{accept/reject}$

security: adv. cannot find  $(S, i, m, \pi)$  s.t.  $m \neq S[i]$  and  
 $\text{verify}(h, i, m, \pi) = \text{accept}$  where  $h = \text{commit}(S)$

# Committing to a list

method 1:  $\text{commit}(S) = h = H(H(m_1), \dots, H(m_n))$

Later: given  $h, m_1$  and  $H(m_2), \dots, H(m_n)$  Bob can check  $S[1] = m_1$

  
proof  $\pi_1$

Problem: long proof!  $(n - 1)$  hash values

Better method: **Merkle tree.** Proof length =  $\log_2 n$  hash values

# Merkle tree

(Merkle 1989)

commitment



$h$

Merkle tree  
commitment

$m_1 \ m_2 \ m_3 \ m_4 \ m_5 \ m_6 \ m_7 \ m_8$



list of values  $S$

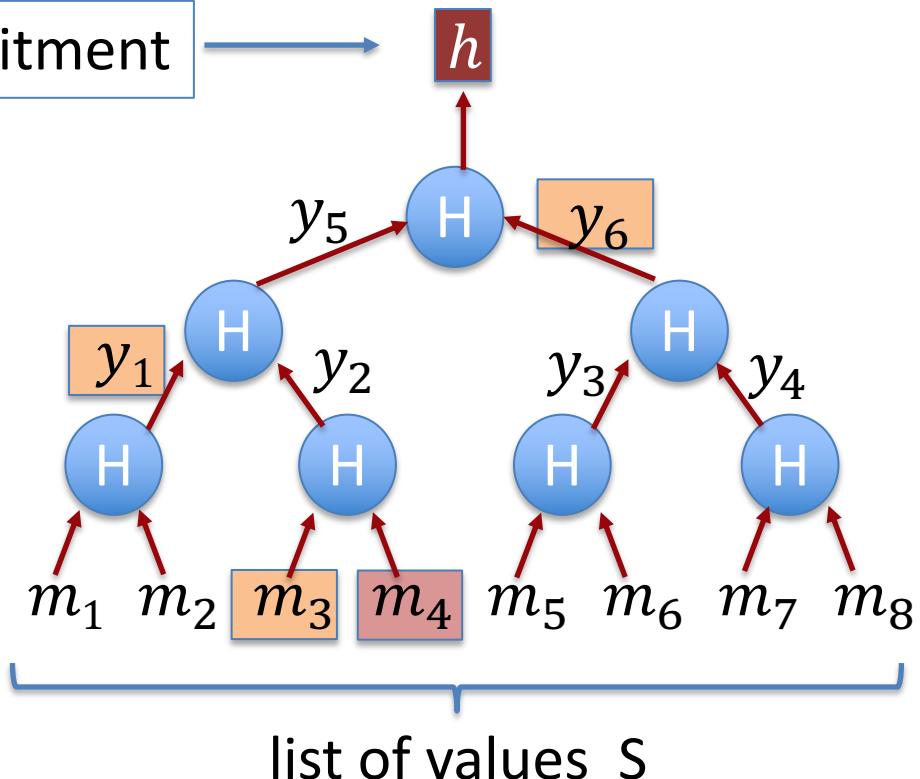
Goal:

- commit to list  $S$
- Later prove  $S[i] = m_i$

# Merkle tree

(Merkle 1989)

commitment



Goal:

- commit to list  $S$
- Later prove  $S[i] = m_i$

To prove  $S[4] = m_4$  ,  
proof  $\pi = (m_3, y_1, y_6)$

length of  $\pi$ :  $\log_2 |S|$

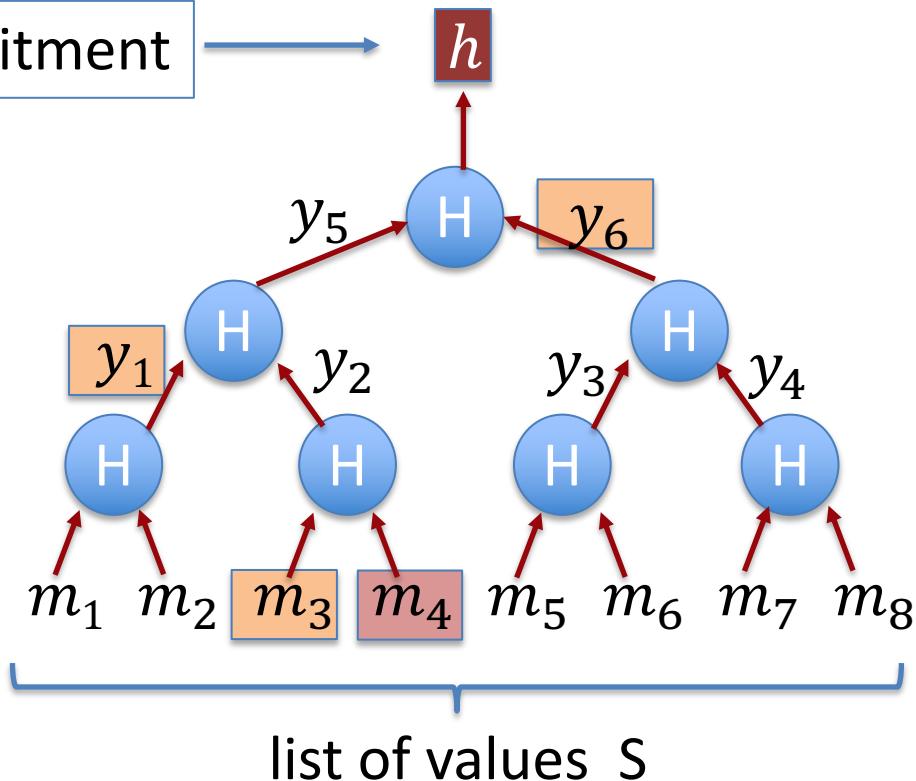
# Merkle tree

(Merkle 1989)

commitment



$h$



To prove  $S[4] = m_4$  ,  
proof  $\pi = (m_3, y_1, y_6)$

Bob does:

$y_2 \leftarrow H(m_3, m_4)$   
 $y_5 \leftarrow H(y_1, y_2)$   
 $h' \leftarrow H(y_5, y_6)$   
accept if  $h = h'$

# Merkle tree

(Merkle 1989)

Thm:  $H \text{ CRHF} \Rightarrow \text{adv. cannot find } (S, i, m, \pi) \text{ s.t. } m \neq S[i] \text{ and}$   
 $\text{verify}(h, i, m, \pi) = \text{accept}$  where  $h = \text{commit}(S)$   
(to prove, prove the contra-positive)

How is this useful? Super useful. Example

- When writing a block of transactions  $S$  to the blockchain, suffices to write  $\text{commit}(S)$  to chain. Keep chain small.
- Later, can prove contents of every Tx.

# Another application: proof of work

**Goal:** computational problem that

- takes time  $\Omega(D)$  to solve, but (D is called the **difficulty**)
- solution takes time  $O(1)$  to verify

How?  $H: X \times Y \rightarrow \{0, 1, 2, \dots, 2^n - 1\}$  e.g.  $n = 256$

- puzzle: input  $x \in X$ , output  $y \in Y$  s.t.  $H(x, y) < 2^n/D$
- verify( $x, y$ ): accept if  $H(x, y) < 2^n/D$

# Another application: proof of work

Thm: if  $H$  is a “random function” then the best algorithm requires  $D$  evaluations of  $H$  in expectation.

Note: this is a parallel algorithm

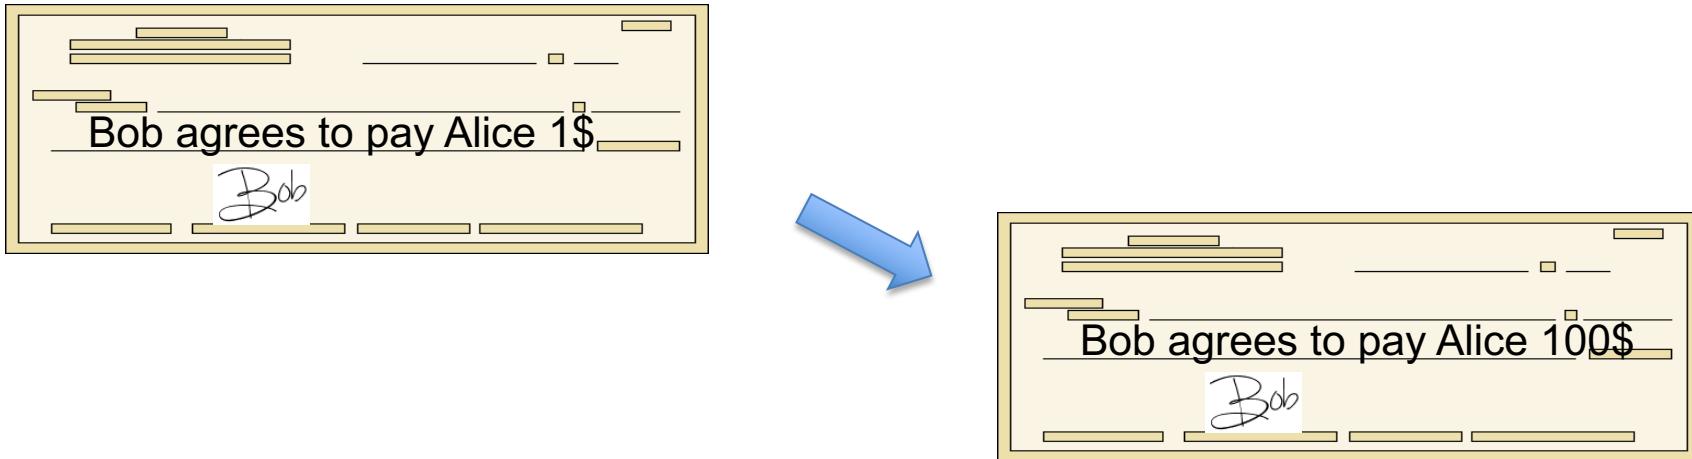
⇒ the more machines I have, the faster I solve the puzzle.

Bitcoin uses  $H(x) = \text{SHA256}(\text{SHA256}(x))$

# Cryptography background: Digital Signatures

# Signatures

Physical signatures: bind transaction to author

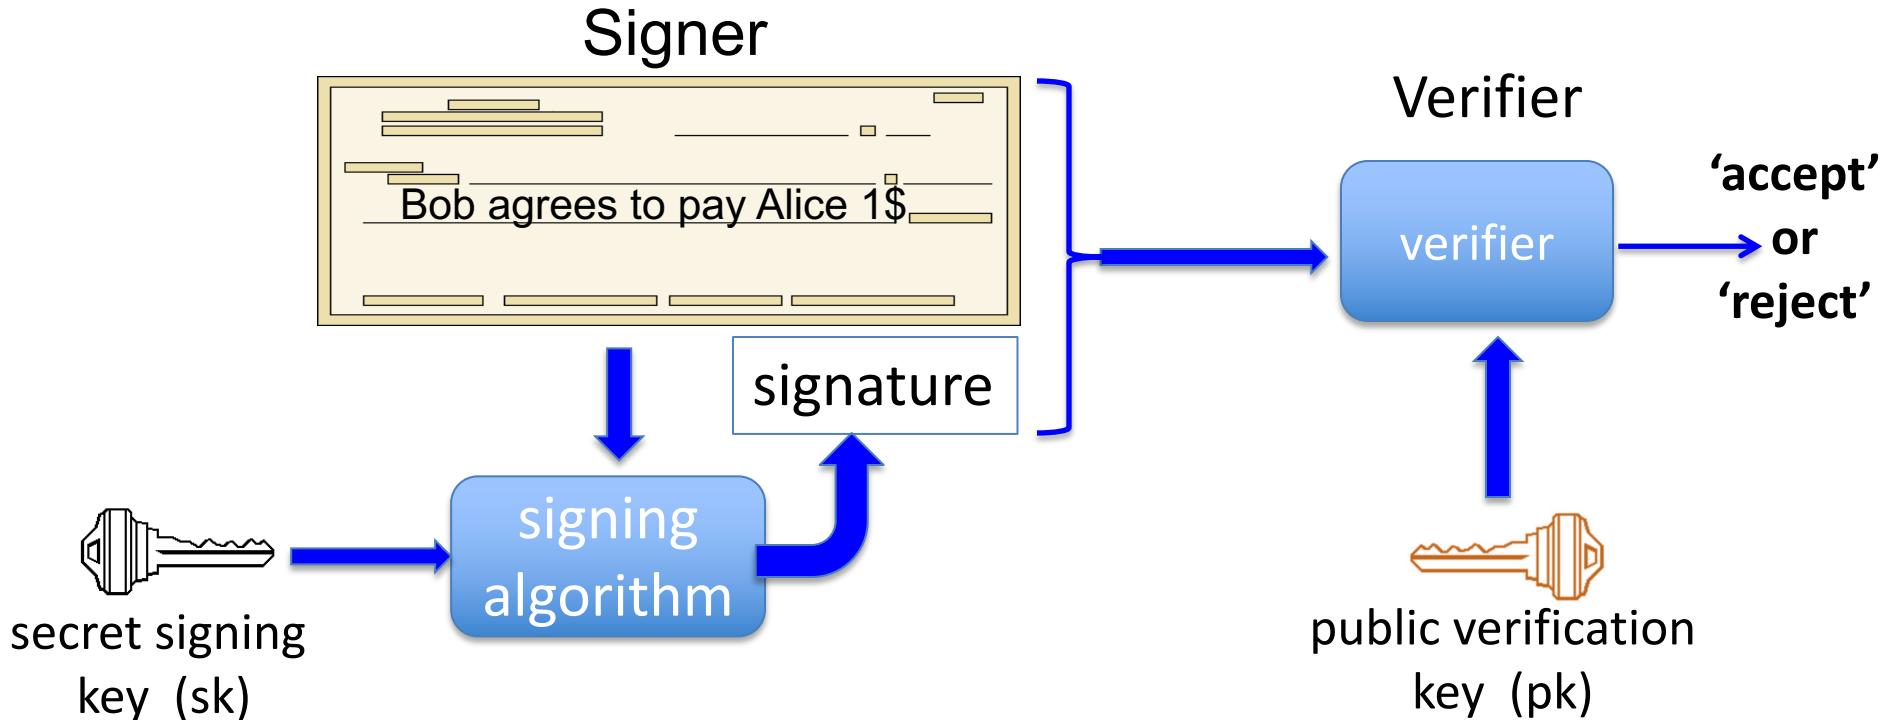


Problem in the digital world:

anyone can copy Bob's signature from one doc to another

# Digital signatures

Solution: make signature depend on document



# Digital signatures: syntax

Def: a signature scheme is a triple of algorithms:

- **Gen()**: outputs a key pair  $(pk, sk)$
- **Sign( $sk, msg$ )** outputs sig.  $\sigma$
- **Verify( $pk, msg, \sigma$ )** outputs ‘accept’ or ‘reject’

Secure signatures: (informal)

Adversary who sees signatures on many messages of his choice, cannot forge a signature on a new message.

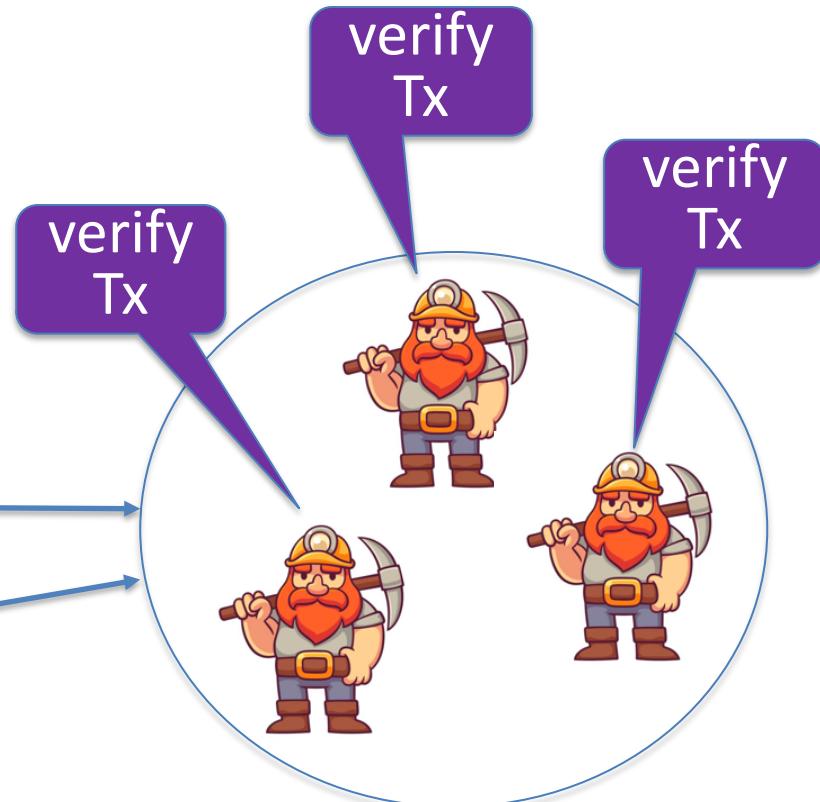
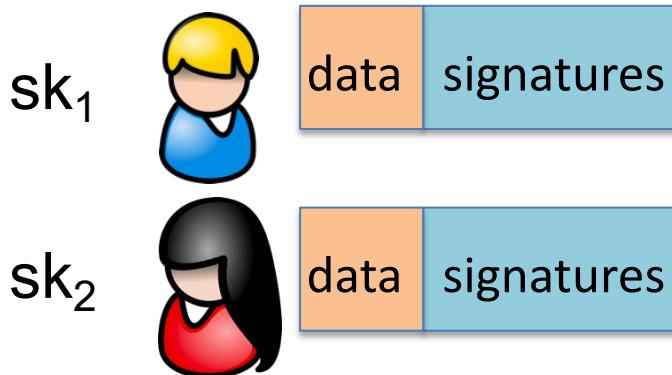
# Families of signature schemes

1. RSA signatures (old ... not used in blockchains):
  - long sigs and public keys ( $\geq 256$  bytes), fast to verify
2. Discrete-log signatures: Schnorr and ECDSA (Bitcoin, Ethereum)
  - short sigs (48 or 64 bytes) and public key (32 bytes)
3. BLS signatures: 48 bytes, aggregatable, easy threshold (Ethereum 2.0, Chia, Dfinity)
4. Post-quantum signatures: long ( $\geq 768$  bytes)

# Signatures on the blockchain

Signatures are used everywhere:

- ensure Tx authorization,
- governance votes,
- consensus protocol votes.



# END OF LECTURE

Next lecture: the Bitcoin blockchain